



Risk-Sensitive Credit Strategy Optimization With Reinforcement Learning

Group 2:

CHANG WEN-YU (01505913)

HONG YANG (01502068)

LI SINUAN (01503010)

PHOO PYAE HSU MYAT (01520414)

ZHANG YINLIANG (01345913)



01

Introduction & Feature Definition





Why This Matters

- Traditional methods rely on **fixed rules** or **income-based policies**
- Fail to adapt to **customer behavior changes**
- Result in **inefficient capital allocation**: Too much risk from **high-risk customers**, Missed revenue from **low-risk customers**

Objective

Develop a **data-driven framework** for **optimizing credit limit decisions**.

- Identify which customers should receive **limit increases**
- Improve **long-term profitability**
- Maintain **acceptable credit risk exposure**

Our Approach

By modeling **repayment patterns** and simulating **limit changes**, we aim to:

- Improve decision quality
- Increase interest revenue
- Reduce losses from defaults

Methodological Overview:



Data Preparation & Feature Engineering


Clean and transform historical credit card usage data into meaningful features for modeling.

Simulator Construction (Classification + Regression)

Build a two-stage supervised ML model to simulate customer behavior.

RL Agent Training

Use DQN in a simulated offline environment.




Policy Evaluation & Backtesting

Compare learned policy against baselines on financial outcomes.

Deployment & Interpretation

Visualize policy decisions and explain feature contributions using SHAP.



Data Source

The dataset is sourced from the **Home Credit Default Risk** competition on **Kaggle**, which provides rich, anonymized credit information at the customer level.

File Name	Description
<code>application_train.csv</code>	Core customer profile and loan application data
<code>credit_card_balance.csv</code>	Monthly credit card usage and payment information
<code>installments_payments.csv</code>	Payment records for past installment loans
<code>POS_CASH_balance.csv</code>	Point-of-sale and cash loan history
<code>bureau.csv</code> + <code>bureau_balance.csv</code>	Third-party credit bureau records and status tracking
<code>previous_application.csv</code>	Records of all prior loan applications and outcomes



Data Wrangling

Features

1

MONTHS_BALANCE(-1 means the freshest balance date)

`-3, -2, -1` as the future data window and `-6,-5,-4` as the prediction window.

2

UR: Utilization Rate (avg over 3 months)

$$UR = \frac{1}{3} \sum_{i=1}^3 \frac{\text{Outstanding Balance}_i}{\text{Credit Limit}_i}$$

3

PR: Payment Rate (avg over 3 months)

$$PR = \frac{1}{3} \sum_{i=1}^3 \frac{\text{AMT_PAYMENT}_i}{\text{AMT_TOTAL_RECEIVABLE}_i}$$

4

TC_i: Total consumer spending in month i (i = 1, 2, 3)

TC₁ = spending in the most recent (last) month

5

EO_i: Operative state in month i: consecutive missed payments

$$EO_1 = \text{non_pay}_1 | EO_2 = \text{non_pay}_2 \cdot (EO_1 + 1) | EO_3 = \text{non_pay}_3 \cdot (EO_2 + 1)$$

6

MP_R: Total number of missed payments in retrospective window

$$MP_R = \sum_{i=1}^3 1 (\text{SK_DPD_DEF}_i > 0)$$



Data Wrangling

Features

7

OB_cday_i: Outstanding balance in month i

$OB_cday_i = \text{AMT_BALANCE at MONTHS_BALANCE} = -i \text{ for } i = 1, 2, 3$

8

P_pday_i: Payment ratio at payment date

$$P_pday_i = \frac{\text{AMT_PAYMENT_TOTAL_CURRENT}_i}{\text{AMT_BALANCE}_i}$$

9

BS: Bureau score proxy

$$BS = \frac{1}{N} \sum_{j=1}^N \text{CREDIT_DAY_OVERDUE}_j$$

10

EI: Estimated Income

$EI = \text{AMT_INCOME_TOTAL}$

11

INT: Interest Rate

Fill NaN with fallback value $\rightarrow 0.1887$

12

N_Months_R: Active months since account opened

$N_Months_R = \text{Count of unique MONTHS_BALANCE entries per customer}$

Data Wrangling

Synthetic Features

13

L_R: Credit Limit in Retrospective Window

Calculation:

$$L_P = \begin{cases} L_R \times \beta & \text{if customer receives an increase} \\ L_R & \text{otherwise} \end{cases}$$

- β : Fixed multiplier (1.5)
- HA_P: A binary flag indicating if $L_P > L_R$

14

L_P: Prospective Credit Limit

Implementation Notes:

- A random sample (20%) of customers can be assigned a limit increase
- Alternatively, the decision can be rule-based (only customers with high payment rate and no missed payments)

15

HA_P: A binary flag indicating if $L_P > L_R$



02

Classify Balance Class Under Supervision Model

Setting Balance Class Threshold

Why Segment by Balance?

Captures Behavioral Structure

- Distinguishes fundamentally different customer types based on balance behavior
- Enables tailored modeling for groups like full payers, light users, and heavy revolvers

Supports Two-Stage Modeling Design

- Classifier pre-filters customers, identifying those likely to carry a balance
- Allows regressors to focus on relevant cases, avoiding noise from zero-balance customers

Benefits

Sharper Provision Estimates

- Focuses credit risk calculations on accounts with actual exposure
- Leads to more accurate and reliable provisioning

Avoids Unnecessary Credit Exposure

- Prevents limit increases for inactive or low-usage customers
- Helps reduce potential losses from unproductive credit

Strengthens RL Policy Outcomes

- Cleaner balance predictions improve reward signal quality
- Helps the RL agent learn safer and more profitable policies

Setting Balance Class Threshold

Distribution Based on AVG_BALANCE

(3-month before decision point average)

Medium=0 50% Half of users carry no balance

75th≈ 81,397 75% Top 25% have balances > 81K

90th≈ 233,371 90% Top 10% carry > 230k

Choose 75K as a threshold, defining top 25% as high usage customers

	Class 2	Class 0	Class 1
Threshold	AVG_BALANCE = 0	$0 < \text{AVG_BALANCE} \leq 75,000$	$\text{AVG_BALANCE} > 75,000$
Description	No balance: Full payers or inactive accounts	Low to moderate balance: Moderate usage, Low Risk	High balance: High usage, Higher risk

Model Training Overview

Use historical customer behavior to **predict future balance behavior class**, enabling smarter credit policy decisions.

Past 3-month behavioral indicators
(MONTHS_BALANCE -7, -8, -9)

Inputs

Labels

Class 0,1,2 based on 3-month
current average balance
(MONTHS_BALANCE -4,-5,-6)

Models Tested
XGBoost,
LightGBM,
Random Forest,
Logistic Regression

Training

Predict


Performance:
Accuracy: 88.43%
Macro F1-score: 85.80%

=== XGBoost ===

	precision	recall	f1-score	support
0	0.80	0.79	0.80	2963
1	0.92	0.92	0.92	7354
accuracy			0.88	10317
macro avg	0.86	0.86	0.86	10317
weighted avg	0.88	0.88	0.88	10317

Model Ranking (on Validation Set):

	Model	Accuracy	Macro F1-score
2	XGBoost	0.884366	0.858081
3	LightGBM	0.884366	0.857642
1	Random Forest	0.882039	0.854869
0	Logistic Regression	0.862557	0.833841



03

Building the RL Environment for Credit Decisions



Transforming Supervised Outputs into RL State:

Preprocessing financial behaviors into state representations for RL training

State Feature Construction:

- 1 Removed redundant features
- 2 Merged 3-month rolling average of UR and PR
- 3 Filled missing values with median
- 4 Binned Δ PROVISION for discrete RL state

Final State Representation:

Utilization Rate

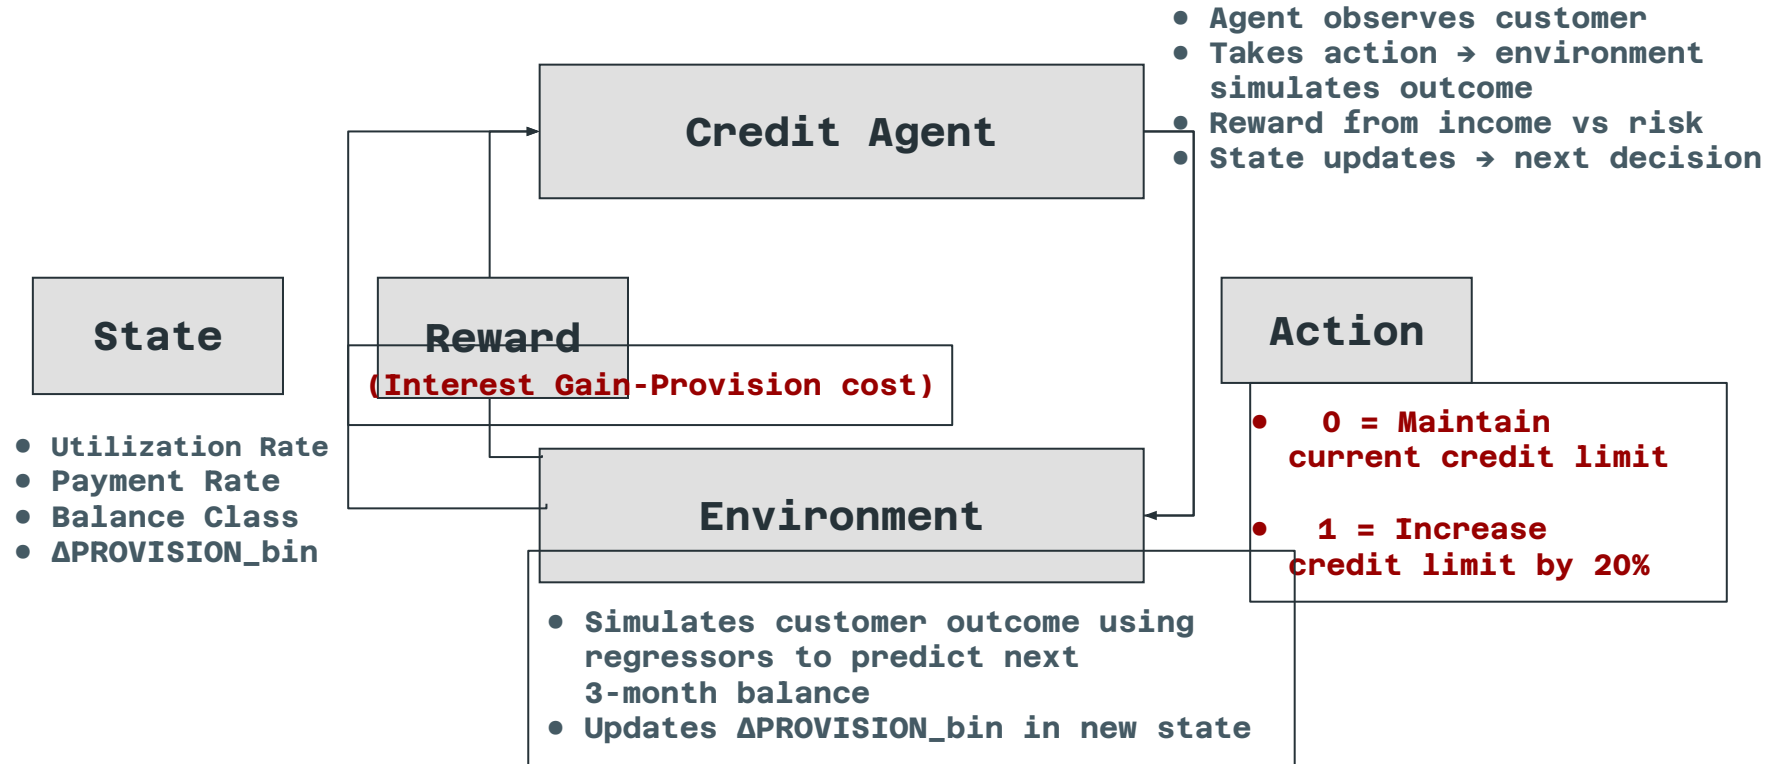
Δ PROVISION_bin

Payment Rate

BALANCE_CLASS = 0,1
(From Phase 1 model)

	PAYMENT_RATE	UTIL_RATE	L_P	HA_P	UR	PR	UR_bin	PR_bin	Δ PROVISION	Δ PROVISION_bin
.	0.579943	0.061619	225000.0	0	0.061619	0.579943	1	11	0.0	49
.	0.662500	0.372428	225000.0	0	0.372428	0.662500	7	13	0.0	49
.	0.066436	0.418133	45000.0	0	0.418133	0.066436	8	1	0.0	49
.	0.106371	0.870598	135000.0	0	0.870598	0.106371	17	2	0.0	49
.	0.032796	0.240578	810000.0	0	0.240578	0.032796	4	0	0.0	49
.
.	0.035868	0.501507	135000.0	0	0.501507	0.035868	10	0	0.0	49
.	0.034624	0.644964	180000.0	0	0.644964	0.034624	12	0	0.0	49
.	0.052997	0.190506	180000.0	0	0.190506	0.052997	3	1	0.0	49
.	0.104444	0.634467	225000.0	0	0.634467	0.104444	12	2	0.0	49
.	0.005809	1.042110	175500.0	0	1.042110	0.005809	0	0	0.0	49

Simulated Environment for Credit Policy Learning:





Reward Function Design: Balancing Profit and Risk

Agent receives a risk-adjusted interest gain - expected loss

Reward Function =

Interest Income

$$3 \times \text{INT} \times \text{BAL} \times (1 - \text{PD})$$

- INT: Customer's interest rate
- BAL: Current balance
- PD: Probability of default
(Assigned by BALANCE_CLASS)

-

Expected Provision Loss

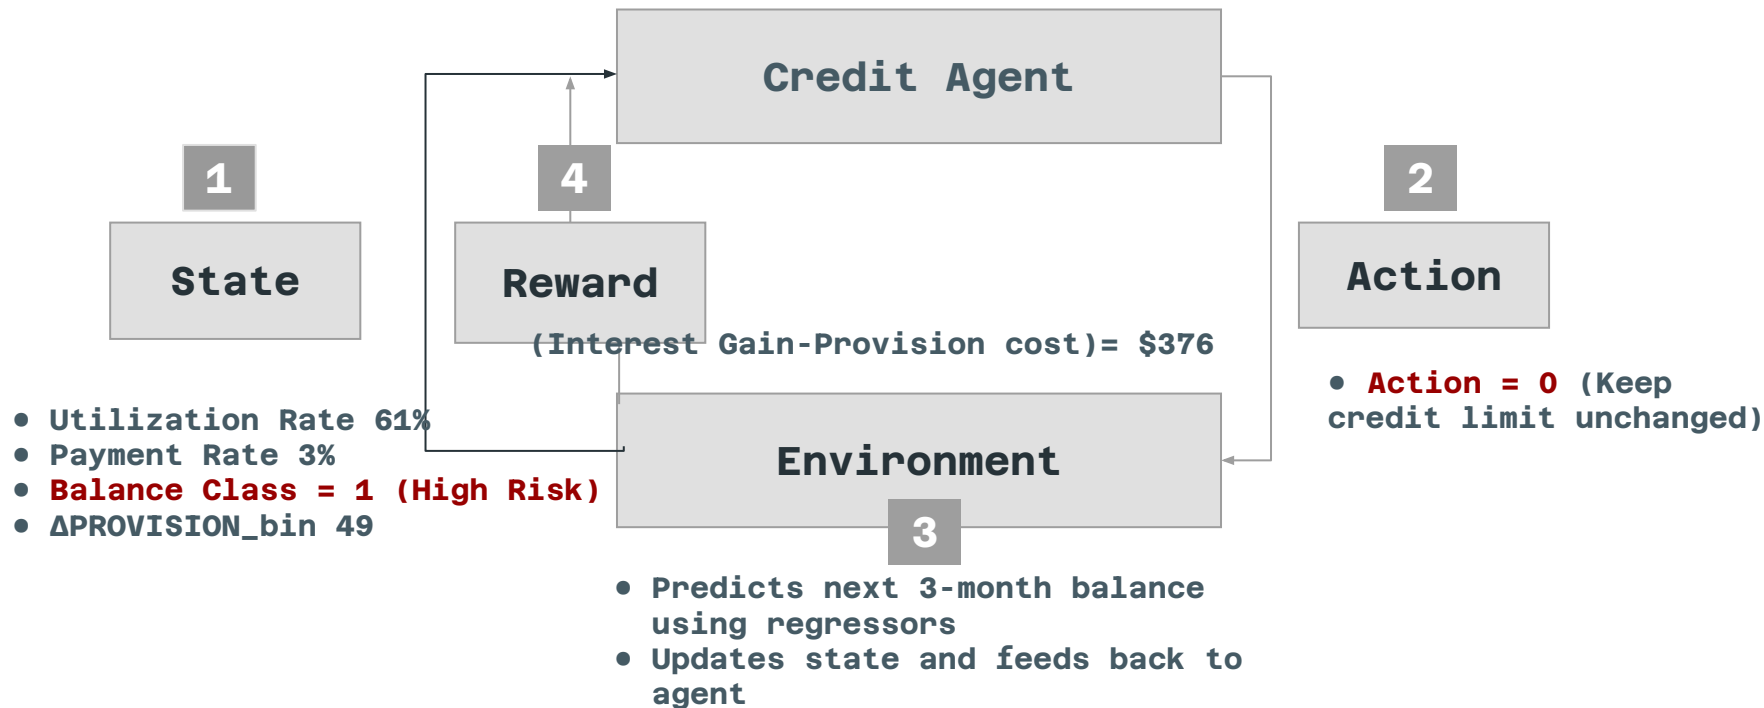
$$\text{PD} \times \text{LGD} \times (\text{BAL}_3 + \text{CCF} \times (\text{L}_P - \text{BAL}_3))$$

- BAL_3: Predicted 3-month balance
- L_P: New credit limit after action
- LGD: Loss Given Default (50%)
- CCF: Credit Conversion Factor (80%)
- L_P - BAL_3: Unused Credit Limit

**Reward guides the agent to balance revenue growth with
risk control in credit policy decisions.**




Simulated Environment Example:



04

Reinforcement Learning - Double Q Network





Why Use DQN for Credit Policy Optimization?




- Handles continuous state space (utilization rate, payment rate, etc.).
- Learns optimal actions via trial and error.
- Balances long-term rewards, not just immediate gains.

Q-Table (Q-learning)

- Typically discrete
- Explicit table $Q(s, a)$
- Uses Bellman equation to update table cells
- Simple, rule-based systems

DQN (Deep Q-Network)

- Can handle continuous or high-dimensional inputs
 - Neural network approximates $Q(s, a)$
 - Minimizes TD error (MSE) through neural networks
 - Large-scale real-world problems
- 

DQN Model Architecture

Network Structure:

- Input: `state_dim = 4`
- Hidden layers: $2 \times \text{Linear}(64)$, ReLU
- Output: `Q-values` for 2 actions

```
def __init__(self, input_dim, output_dim):  
    super(DQN, self).__init__()  
    self.fc1 = nn.Linear(input_dim, 64)  
    self.fc2 = nn.Linear(64, 64)  
    self.out = nn.Linear(64, output_dim)
```

- ➡ **Lightweight and fast for convergence**
- ➡ **Sufficiently expressive for a small discrete action space**



DQN Key Hyperparameters




Parameter	Value	Purpose
Episodes	100	Training cycles
γ (Gamma)	0.9	Discount factor for future rewards
ϵ -decay	0.995 \rightarrow 0.05	Epsilon-greedy exploration
Batch size	64	Sample size for replay buffer
Target Update	every 10 episodes	Stabilize Q-learning
Optimizer	Adam	Learning Q-network



Training Result: Reward Curve

- High volatility across 200-episode
- Upper reward envelope increases gradually
- High-reward episodes become more frequent over time
- Indicates partial learning progress and agent's ability to exploit favorable cases



Reward Curve Instability Analysis

Possible Causes:

Environment stochasticity - Customer sequence and response vary

Volatile reward design - Rewards depend on predicted balances, INT, PD

Cold Start Instability - Early Q-network is untrained, transitions still limited

Possible Solutions:

Shuffle customer order each episode

Train over more episodes to average out randomness

Clip or normalize rewards to reduce extreme value swings

Modularize reward components (e.g., separate interest and provision)

Warm up replay buffer with 2-3× batch size before training starts

Use running/moving average to smooth reward tracking




Action Distribution Analysis



- **Action 0:** 1,040,609 times ($\approx 55.3\%$) (Keep credit limit unchanged)
- **Action 1:** 840,969 times ($\approx 44.7\%$)
- The agent shows a slight preference for Action 0
- Suggests Action 0 may be slightly more rewarding or safer

Policy Evaluation



Benchmark Comparison


- Maintain Limit
- Always Increase
- Random

DQN RL Strategy

- Achieved **highest average reward**
- **Significantly outperformed** all baseline strategies
- Demonstrated ability to learn an **effective dynamic credit limit policy**

05

Conclusion and Future Directions



Conclusion & Future Directions

Key Takeaways

- Successfully implemented a functioning **DQN-based credit limit optimizer**.
- Balanced **revenue and credit risk** using simulated customer behavior and financial outcomes.
- Demonstrated proof of concept for **data-driven decision support** in consumer credit portfolios.

Major Future Improvements

- **Personalized Default Prediction:** Replace fixed PD with customer-level probability of default models.
- **Markov-Based Transitions:** Model behavior evolution over time using **state transition dynamics**.
- **Multi-Period Optimization:** Track cumulative rewards to reflect long-term profitability.
- **Feature Enrichment:** Integrate **repayment trends, RFM scores, and category-level spending**.

Thank You !

